

spreadsheet sorts, it moves the associated record (the satellite data) with the key. When describing a sorting algorithm, we focus on the keys, but it is important to remember that there usually is associated satellite data.


In this book, we'll typically describe algorithms as procedures written in a *pseudocode* that is similar in many respects to C, C++, Java, Python,<sup>1</sup> or JavaScript. (Apologies if we've omitted your favorite programming language. We can't list them all.) If you have been introduced to any of these languages, you should have little trouble understanding algorithms "coded" in pseudocode. What separates pseudocode from real code is that in pseudocode, we employ whatever expressive method is most clear and concise to specify a given algorithm. Sometimes the clearest method is English, so do not be surprised if you come across an English phrase or sentence embedded within a section that looks more like real code. Another difference between pseudocode and real code is that pseudocode often ignores aspects of software engineering—such as data abstraction, modularity, and error handling—in order to convey the essence of the algorithm more concisely.

We start with *insertion sort*, which is an efficient algorithm for sorting a small number of elements. Insertion sort works the way you might sort a hand of playing cards. Start with an empty left hand and the cards in a pile on the table. Pick up the first card in the pile and hold it with your left hand. Then, with your right hand, remove one card at a time from the pile, and insert it into the correct position in your left hand. As Figure 2.1 illustrates, you find the correct position for a card by comparing it with each of the cards already in your left hand, starting at the right and moving left. As soon as you see a card in your left hand whose value is less than or equal to the card you're holding in your right hand, insert the card that you're holding in your right hand just to the right of this card in your left hand. If all the cards in your left hand have values greater than the card in your right hand, then place this card as the leftmost card in your left hand. At all times, the cards held in your left hand are sorted, and these cards were originally the top cards of the pile on the table.

The pseudocode for insertion sort is given as the procedure INSERTION-SORT on the facing page. It takes two parameters: an array  $A$  containing the values to be sorted and the number  $n$  of values of sort. The values occupy positions  $A[1]$  through  $A[n]$  of the array, which we denote by  $A[1:n]$ . When the INSERTION-SORT procedure is finished, array  $A[1:n]$  contains the original values, but in sorted order.

---

<sup>1</sup> If you're familiar with only Python, you can think of arrays as similar to Python lists.



**Figure 2.1** Sorting a hand of cards using insertion sort.


```

INSERTION-SORT( $A, n$ )
1  for  $i = 2$  to  $n$ 
2       $key = A[i]$ 
3      // Insert  $A[i]$  into the sorted subarray  $A[1 : i - 1]$ .
4       $j = i - 1$ 
5      while  $j > 0$  and  $A[j] > key$ 
6           $A[j + 1] = A[j]$ 
7           $j = j - 1$ 
8       $A[j + 1] = key$ 

```

### Loop invariants and the correctness of insertion sort

Figure 2.2 shows how this algorithm works for an array  $A$  that starts out with the sequence  $\langle 5, 2, 4, 6, 1, 3 \rangle$ . The index  $i$  indicates the “current card” being inserted into the hand. At the beginning of each iteration of the **for** loop, which is indexed by  $i$ , the *subarray* (a contiguous portion of the array) consisting of elements  $A[1 : i - 1]$  (that is,  $A[1]$  through  $A[i - 1]$ ) constitutes the currently sorted hand, and the remaining subarray  $A[i + 1 : n]$  (elements  $A[i + 1]$  through  $A[n]$ ) corresponds to the pile of cards still on the table. In fact, elements  $A[1 : i - 1]$  are the elements *originally* in positions 1 through  $i - 1$ , but now in sorted order. We state these properties of  $A[1 : i - 1]$  formally as a *loop invariant*:



**Figure 2.2** The operation of  $\text{INSERTION-SORT}(A, n)$ , where  $A$  initially contains the sequence  $\langle 5, 2, 4, 6, 1, 3 \rangle$  and  $n = 6$ . Array indices appear above the rectangles, and values stored in the array positions appear within the rectangles. **(a)–(e)** The iterations of the **for** loop of lines 1–8. In each iteration, the blue rectangle holds the key taken from  $A[i]$ , which is compared with the values in tan rectangles to its left in the test of line 5. Orange arrows show array values moved one position to the right in line 6, and blue arrows indicate where the key moves to in line 8. **(f)** The final sorted array.

At the start of each iteration of the **for** loop of lines 1–8, the subarray  $A[1 : i - 1]$  consists of the elements originally in  $A[1 : i - 1]$ , but in sorted order.

Loop invariants help us understand why an algorithm is correct. When you're using a loop invariant, you need to show three things:

**Initialization:** It is true prior to the first iteration of the loop.

**Maintenance:** If it is true before an iteration of the loop, it remains true before the next iteration.

**Termination:** The loop terminates, and when it terminates, the invariant—usually along with the reason that the loop terminated—gives us a useful property that helps show that the algorithm is correct.

When the first two properties hold, the loop invariant is true prior to every iteration of the loop. (Of course, you are free to use established facts other than the loop invariant itself to prove that the loop invariant remains true before each iteration.)

A loop-invariant proof is a form of mathematical induction, where to prove that a property holds, you prove a base case and an inductive step. Here, showing that the invariant holds before the first iteration corresponds to the base case, and showing that the invariant holds from iteration to iteration corresponds to the inductive step.

The third property is perhaps the most important one, since you are using the loop invariant to show correctness. Typically, you use the loop invariant along with the condition that caused the loop to terminate. Mathematical induction typically applies the inductive step infinitely, but in a loop invariant the “induction” stops when the loop terminates.

Let's see how these properties hold for insertion sort.

**Initialization:** We start by showing that the loop invariant holds before the first loop iteration, when  $i = 2$ .<sup>2</sup> The subarray  $A[1:i-1]$  consists of just the single element  $A[1]$ , which is in fact the original element in  $A[1]$ . Moreover, this subarray is sorted (after all, how could a subarray with just one value not be sorted?), which shows that the loop invariant holds prior to the first iteration of the loop.

**Maintenance:** Next, we tackle the second property: showing that each iteration maintains the loop invariant. Informally, the body of the **for** loop works by moving the values in  $A[i-1]$ ,  $A[i-2]$ ,  $A[i-3]$ , and so on by one position to the right until it finds the proper position for  $A[i]$  (lines 4–7), at which point it inserts the value of  $A[i]$  (line 8). The subarray  $A[1:i]$  then consists of the elements originally in  $A[1:i]$ , but in sorted order. *Incrementing*  $i$  (increasing its value by 1) for the next iteration of the **for** loop then preserves the loop invariant.

A more formal treatment of the second property would require us to state and show a loop invariant for the **while** loop of lines 5–7. Let's not get bogged down in such formalism just yet. Instead, we'll rely on our informal analysis to show that the second property holds for the outer loop.

**Termination:** Finally, we examine loop termination. The loop variable  $i$  starts at 2 and increases by 1 in each iteration. Once  $i$ 's value exceeds  $n$  in line 1, the loop terminates. That is, the loop terminates once  $i$  equals  $n + 1$ . Substituting  $n + 1$  for  $i$  in the wording of the loop invariant yields that the subarray  $A[1:n]$  consists of the elements originally in  $A[1:n]$ , but in sorted order. Hence, the algorithm is correct.

This method of loop invariants is used to show correctness in various places throughout this book.

### Pseudocode conventions

We use the following conventions in our pseudocode.

- Indentation indicates block structure. For example, the body of the **for** loop that begins on line 1 consists of lines 2–8, and the body of the **while** loop that

---

<sup>2</sup> When the loop is a **for** loop, the loop-invariant check just prior to the first iteration occurs immediately after the initial assignment to the loop-counter variable and just before the first test in the loop header. In the case of INSERTION-SORT, this time is after assigning 2 to the variable  $i$  but before the first test of whether  $i \leq n$ .